

# Dialogue Is Not Enough to Make a Communicative BabyLM (But Neither Is Developmentally Inspired Reinforcement Learning)



Francesca Padovani<sup>1\*</sup> & Bastian Bunzeck<sup>2\*</sup> & Manar Ali<sup>2</sup> & Omar Momen<sup>2</sup> & Arianna Bisazza<sup>1</sup> & Hendrik Buschmeier<sup>2</sup> & Sina Zarrieß<sup>2</sup>  
| CLCG, University of Groningen | CRC 1646, Bielefeld University [f.padovani@rug.nl](mailto:f.padovani@rug.nl), [bastian.bunzeck@uni-bielefeld.de](mailto:bastian.bunzeck@uni-bielefeld.de)

**General Idea:** simulating more fluent and contextually plausible spoken interactions between child and caregivers

## LLAMALOGUE

(model trained on dialogue data, triplets of three consecutive turns)

\*MOT: oh yes .

\*CHI: make it flat .

\*MOT: okay .

\*MOT: what're you making ?

\*CHI: a house .

\*MOT: oh very nice .

\*MOT: what do you think that is ?

\*CHI: eggs .

\*MOT: eggs

### AUTOREGRESSIVE MODELING

Llama model trained on 10M of CHILDES tokens for 10 epochs.

## DPO (Direct Preference Optimization)

### 1. Naturalistic data

Prompt (MOT)	Appropriate (CHI)	Random (CHI)
what is that ?	it looks like a gun .	you do it like that .
pull the string .	and where do they hook it ?	do you know what it was ?
I think they just throw it on the side .	you know what Mom ?	I get this hole .
what are you playing with huh toys ?	there's a dog .	there's the sports .
the bottom ones come off .	want to know what ?	we stole the brush .
can you say that ?	okay the hungry hungry caterpillar .	yeah I want that too .
what is it ?	a baby caterpillar !	I'm just pretending .
what is it ?	I don't .. know !	put my dress down .

### 2. Synthetic data

Prompt (MOT)	Appropriate (CHI)	Random (CHI)
what is that ?	it is a ball.	you do it like that .
pull the string .	i can do it!	do you know what it was ?
I think they just throw it on the side .	the garbage truck.	I get this hole .
what are you playing with huh toys ?	i'm playing with this toy, it's a nice toy.	there's the sports .
the bottom ones come off .	the bottom ones come off .	we stole the brush .
can you say that ?	yes, mother can say that.	yeah I want that too .
what is it ?	i don't know.	I'm just pretending .
what is it ?	a big tree.	put my dress down .


## PPO (Proximal Policy Optimization)

### What constitutes a "successful" communicative turn?

Experimenting with four reward functions:

- ▶ 1-gram BLEU Reward
- ▶ Semantic Similarity Reward
- ▶ A quality score derived from an LLM (Llama-3.2-3B)
- ▶ An uncertainty-based measuring LLM confidence in processing child responses

A teacher LLM simulates short child-like answer that shows understanding of the caregiver utterance.



## Prompt to Llama-3.2-3B

▶ Reward functions are computed by comparing the teacher-generated responses to the output produced by *llamalogue* when prompted with the same utterance.

▶ Also the synthetic sentences used in the DPO dataset are generated with the same prompt.

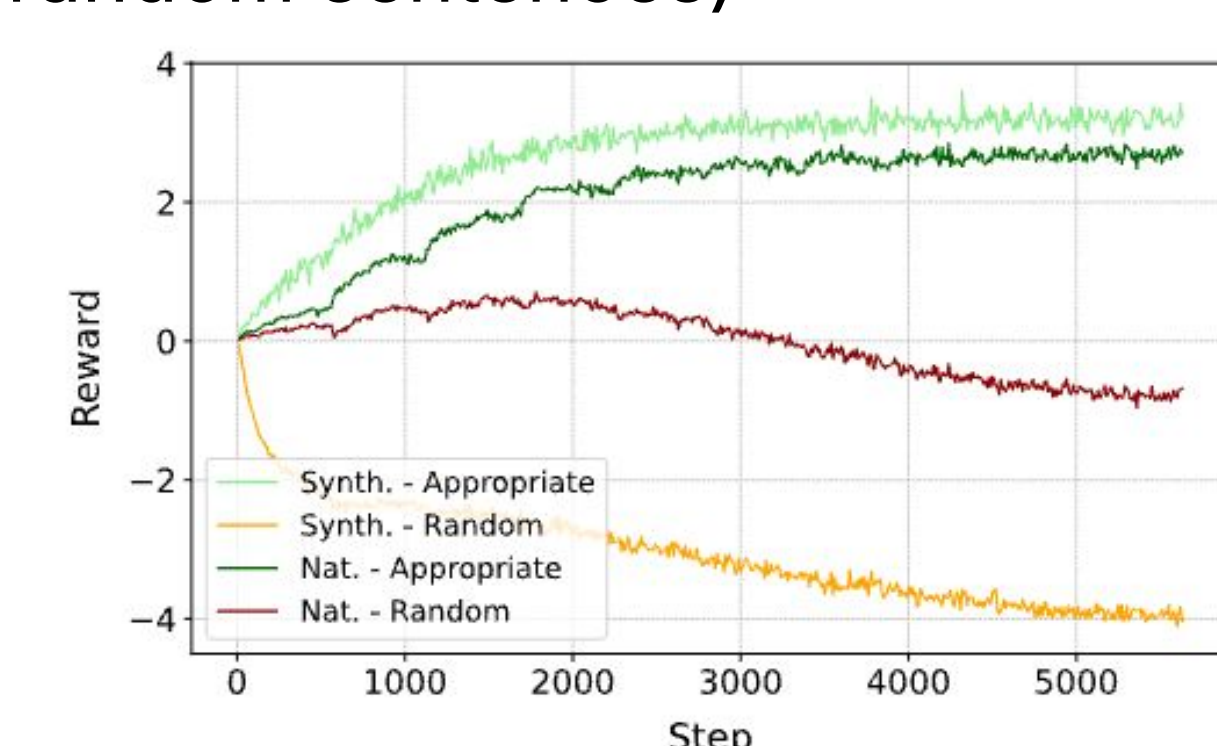
You are a young child having a conversation with your mother.  
When your mother says something, you should answer as a typical and natural-sounding child. Do NOT repeat her words. Instead, give a new, relevant answer that shows understanding.  
Keep it short and child-like.

\*MOT: I think they just throw it on the side .

\*CHI:


## DPO Finetuning

Reward trends (for the appropriate and random sentences)




## PPO Finetuning


▶ **nltk** to compute a smoothed unigram BLEU score (BLEU-1) between *llamalogue*'s response and teacher reference answer



▶ **all-MiniLM-L6-v2** model from SentenceTransformers to compute the cosine similarity

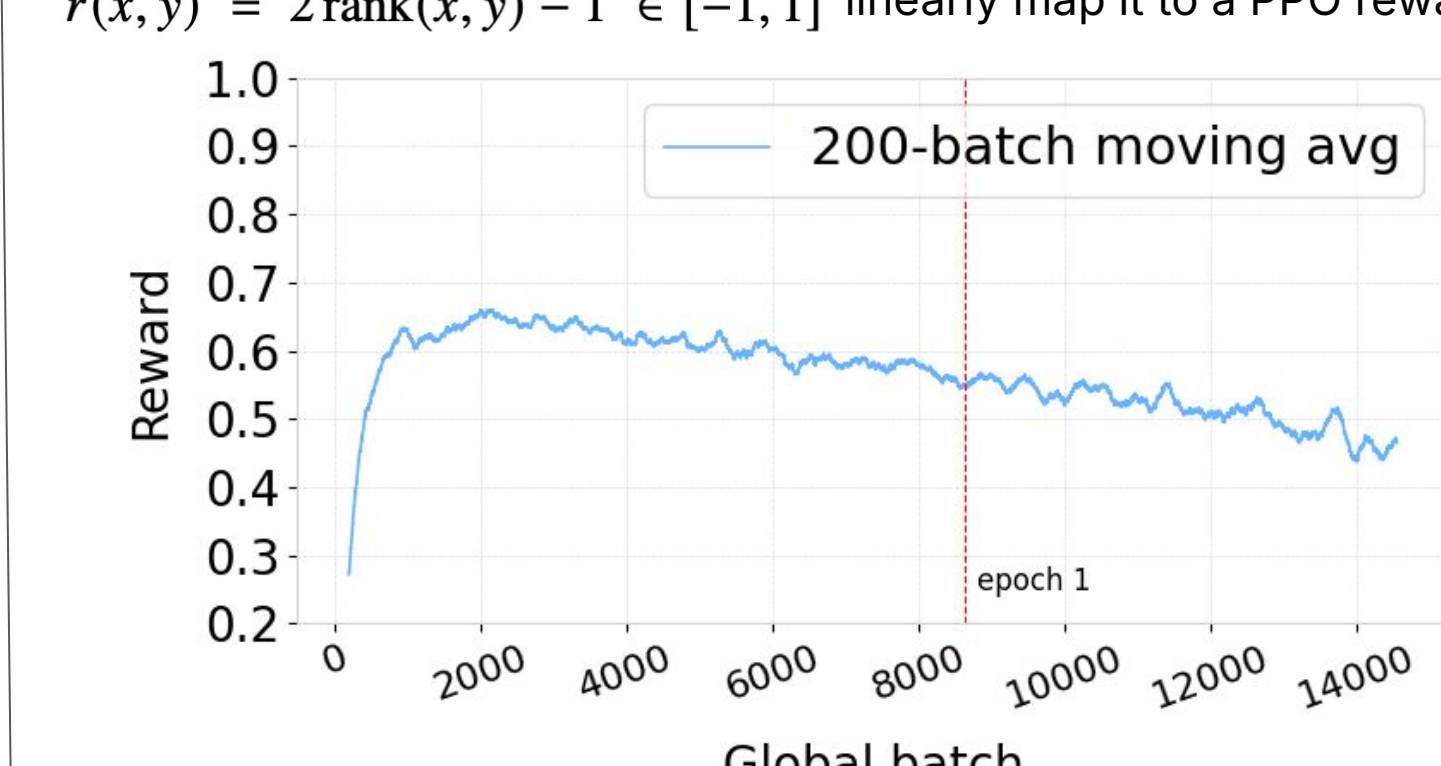


▶ **Olmo LLM** generate numerical quality **score** (from 0 to 5) based on contextual appropriateness of child response



▶ **Confidence based reward**

$\ell_{\text{baby}} = \log P_{\text{teacher}}(\bar{y} | x)$  log-likelihood for 10 LLM-answers  
 $\text{rank}(x, \bar{y}) = \frac{1}{10} \sum_{i=1}^{10} \mathbb{1}\{\ell_i \leq \ell_{\text{baby}}\} \in [0, 1]$  compute normalized rank  
 $r(x, \bar{y}) = 2 \text{rank}(x, \bar{y}) - 1 \in [-1, 1]$  linearly map it to a PPO reward



## Evaluation

**BabyLM Challenge evaluation suite:** both zero-shot and fine-tuning

**Custom Benchmarks:** dialogue minimal pairs, Zorro and lexical decision task.

	Task	llamalogue	DPO		PPO				Baseline
			Natural.	Synth.	Bleu	SemSim	LLM Score	Conf.	
Zero-shot (Baby LM)	BLiMP	56.05	55.64	55.51	55.14	<b>56.36</b>	55.31	55.10	72.16
	BLiMP suppl.	51.06	49.97	<b>51.67</b>	51.33	51.48	50.58	49.45	61.22
	COMPS	51.62	51.51	<b>51.63</b>	50.66	51.58	51.25	51.59	—
	Entity tracking	30.66	32.66	31.29	16.20	34.64	<b>36.03</b>	34.05	28.06
	EWoK	50.19	50.12	<b>50.82</b>	49.65	49.62	50.12	50.81	51.92
	Read. (eye track.)	<b>3.88</b>	3.57	1.16	3.43	2.85	3.73	3.35	9.08
	Read. (self-paced)	1.43	1.35	0.44	<b>1.99</b>	1.04	1.30	1.14	3.5
	Wug adj.	0.45	0.52	0.16	0.13	0.01	<b>0.55</b>	0.41	38.5
	Wug past	-0.03	-0.01	-0.05	-0.15	-0.18	<b>-0.01</b>	-0.19	—
	AoA	-79.6	0	0	-80.1	0	-76.6	-78.7	—
FT (Add'l)	(Super)GLUE	51.82	<b>51.72</b>	51.77	51.12	52.10	51.69	<b>51.92</b>	67.91
	Lexical decision	40.3	<b>40.5</b>	<b>41.3</b>	40.7	39.7	40.2	40.8	57.2
	Zorro	<b>65.5</b>	<b>64.8</b>	62.7	62.5	64.7	65.2	63.7	77.7
	Dia. MP (Words)	64.3	<b>68.4</b>	64.9	62	61.1	60.6	63.7	58.1
	Dia. MP (Tokens)	63.8	<b>67.6</b>	64.3	61	63.6	62.5	62.4	57.9


## Results

### BabyLM Challenge evaluation suite

1. **Llamalogue** underperforms the BabyLM interaction baseline on most formal benchmarks, **except for entity tracking**.
2. It scores lower on **BLiMP** and **AoA prediction**, with modest results on **EWoK, COMPS, (Super)GLUE**, and **wug tests** (compared to other small-track submissions).

### Custom Benchmarks

1. On dialogue minimal pairs task - that aligns with the pre-training goal of *llamalogue* - it exhibits a clear advantage over the baseline.
2. Our model achieves accuracy of 65.5% on Zorro, but is outperformed by the baseline.
3. Llamalogue also falls behind the interactive baseline on the lexical decision task.



"We find that pre-training small LMs on child-caregiver dialogues somehow preserves formal linguistic competence and that Direct Preference alignment (DPO) better captures conversational intent than PPO (Proximal Policy Optimization), but neither improves formal language skills."